

code

August 17, 2023

0.1 Assignment 1

0.1.1 Name: Bannuru Rohit Kumar Reddy

0.1.2 Roll Number: 21CS30011

Libraries :

```
[150]: # import all the necessary libraries here :

import pandas as pd

import numpy as np
from sklearn.model_selection import train_test_split

from sklearn.metrics import accuracy_score, precision_score, recall_score
```

Loading the dataset and analysing the type of data

```
[151]: df = pd.read_excel('../..//dataset/logistic-regression/Pumpkin_Seeds_Dataset.
↳xlsx')
print(df.shape)
print(df.head(1))
```

(2500, 13)

	Area	Perimeter	Major_Axis_Length	Minor_Axis_Length	Convex_Area	\
0	56276	888.242	326.1485	220.2388	56831	
	Equiv_Diameter	Eccentricity	Solidity	Extent	Roundness	Aspect_Ration \
0	267.6805	0.7376	0.9902	0.7453	0.8963	1.4809
	Compactness	Class				
0	0.8207	Çerçvelik				

PreProcessing the Data:

```
[152]: # PreProcessing the data :

# Encode the 'Class' column to numerical values
class_mapping = {'Çerçvelik': 0, 'Ürgüp Sivrisi': 1}
df['Class'] = df['Class'].map(class_mapping)
```

```

# Split the dataset into train, validation, and test sets
train_data, temp_data = train_test_split(df, test_size=0.5, random_state=42)
valid_data, test_data = train_test_split(temp_data, test_size=0.4,
    random_state=42)

# Separate features and target variables
X_train = train_data.drop('Class', axis=1).values
y_train = train_data['Class'].values

X_valid = valid_data.drop('Class', axis=1).values
y_valid = valid_data['Class'].values

X_test = test_data.drop('Class', axis=1).values
y_test = test_data['Class'].values

# Normalize/Standardize the features
mean = np.mean(X_train, axis=0)
std = np.std(X_train, axis=0)
X_train = (X_train - mean) / std
X_valid = (X_valid - mean) / std
X_test = (X_test - mean) / std

# Add a column of 1's (bias term) to X_train, X_valid, and X_test using NumPy
bias_column_train = np.ones((X_train.shape[0], 1))

# Concatenate the bias column with the original X_train, X_valid, and X_test
X_train = np.hstack((bias_column_train, X_train))
X_valid = np.hstack((np.ones((X_valid.shape[0], 1)), X_valid))
X_test = np.hstack((np.ones((X_test.shape[0], 1)), X_test))

print(X_train)
print(y_train)

```

```

[[ 1.          0.60882926  0.6459801  ... -0.30180065  0.01488933
  -0.14217799]
 [ 1.          0.32818862 -0.24722132 ...  1.58513179 -1.47631136
   1.68881593]
 [ 1.          1.2310118   1.65276095 ... -1.43895861  1.61791432
  -1.50887079]
 ...
 [ 1.          1.01272767  0.49924928 ...  0.97281596 -0.95309963
   0.98300334]
 [ 1.          0.07163201 -0.07652727 ...  0.42476652 -0.42705291
   0.34458931]
 [ 1.         -0.105833   -0.10490407 ...  0.06773106 -0.05661523

```

```
-0.03920798]]  
[1 0 1 ... 0 0 0]
```

Logistic Regression :

```
[153]: # Defining the model :  
  
def sigmoid(z):  
    return 1 / (1 + np.exp(-z))  
  
def predict(X, weights):  
    return sigmoid(np.dot(X, weights.T))  
  
def logistic_regression(X, y, num_epochs, learning_rate):  
    num_samples, num_features = X.shape  
    weights = np.zeros(num_features)  
  
    for epoch in range(num_epochs):  
        y_pred = predict(X, weights)  
        gradient = np.dot(X.T, (y_pred - y)) / num_samples  
        weights -= learning_rate * gradient  
  
    return weights
```

Predicting values for test data based :

```
[154]: # Training the logistic regression model  
num_epochs = 1000  
learning_rate = 0.001  
weights = logistic_regression(X_train, y_train, num_epochs, learning_rate)  
  
# Make predictions on the validation set  
y_pred_valid = predict(X_valid, weights)  
y_pred_valid_class = np.round(y_pred_valid)  
  
# Make predictions on the test set  
y_pred_test = predict(X_test, weights)  
y_pred_test_class = np.round(y_pred_test)
```

Evaluating how well the model is performing on training and test data :

```
[155]: def evaluate_metrics(y_true, y_pred):  
    accuracy = accuracy_score(y_true, y_pred)  
    precision = precision_score(y_true, y_pred)  
    recall = recall_score(y_true, y_pred)  
    return accuracy, precision, recall  
  
# Evaluate on validation set
```

```

accuracy_valid, precision_valid, recall_valid = evaluate_metrics(y_valid,
    ↪y_pred_valid_class)

# Evaluate on test set
accuracy_test, precision_test, recall_test = evaluate_metrics(y_test,
    ↪y_pred_test_class)

print("Validation Set Metrics:")
print(f"Accuracy: {accuracy_valid:.5f}")
print(f"Precision: {precision_valid:.5f}")
print(f"Recall: {recall_valid:.5f}")

print("\nTest Set Metrics:")
print(f"Accuracy: {accuracy_test:.5f}")
print(f"Precision: {precision_test:.5f}")
print(f"Recall: {recall_test:.5f}")

```

Validation Set Metrics:

Accuracy: 0.84533

Precision: 0.85797

Recall: 0.81543

Test Set Metrics:

Accuracy: 0.85400

Precision: 0.83750

Recall: 0.85532

[]: