




Algorithms – I (CS29003/203)

Autumn 2022, IIT Kharagpur

Priority Queues and Heaps



Breakthrough during the Vacation





Resources

- Apart from the book
- UC Davis ECS 36C Course by Prof. Joël Porquet-Lupine



Introduction

- Let us think about a scenario where you are performing the following tasks simultaneously
 - Start a long code compilation
 - Refresh moodle page
 - Send Gmail chat message
 - Receive email notification
- Objective: Execute multiple processes until completion, but keep the system responsive



- Are processes executed until completion?
- Will the system feel responsive?




Process Scheduling

- Round-robin scheduling
 - Add processes to a (typically FIFO) queue
 - Execute them for equal chunks of time



- Will the system feel more responsive? Can we do better?
- Priority scheduling
 - Associate priority with each process, e.g., short processes get higher priority
 - Execute the processes with the highest priority first

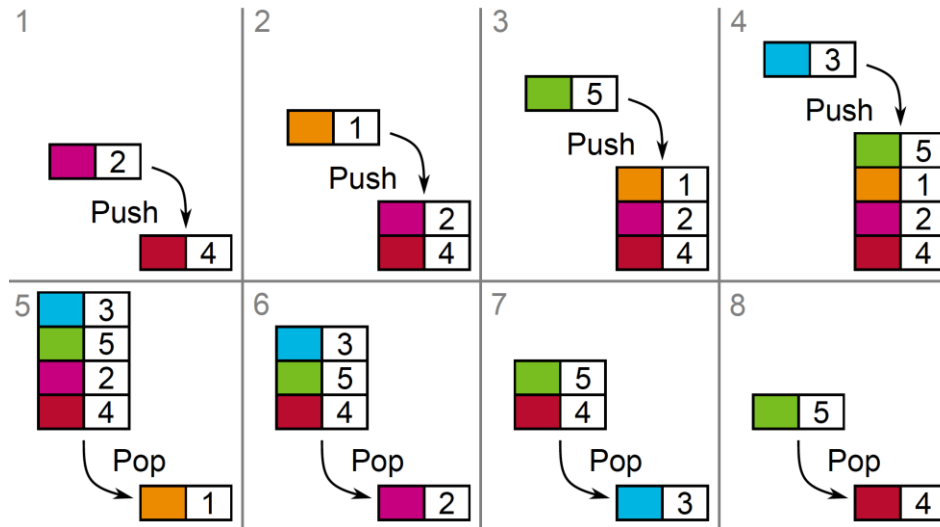


A very good resource on process scheduling:
[Prof. Krzyzanowski's lecture notes](#)



Priority Queue



- **Definition:** **Queue** where each item is associated to a priority (i.e., a comparable key)
 - **Push/Insert:** add item and associated priority
 - **Pop/Pull:** remove item with highest priority





Naïve Implementations

- Sorted list
 - Keep sorted by priority when pushing
 - Pop highest priority item from front
- Binary Search Tree (self-balanced)
 - Restructure tree by priority when inserting
 - Return highest priority item from minimum of tree





List Implementations

- Struct node

```
// Node
struct node {
    // Lower values indicate higher priority
    int priority;
    struct node* next;
};
```

- Function to create a new node

```
// Function to Create A New Node
struct node* newNode(int p)
{
    struct node* temp = (struct node*)malloc(sizeof(struct node));
    temp->priority = p;
    temp->next = NULL;
    return temp;
}
```

- Function to check if queue is empty

```
bool isEmpty(struct node* head)
{
    return head == NULL;
}
```




List Implementations

- Get value of the maximum priority (minimum element)

```
int getMin(struct node* head)
{
    if(isEmpty(head))
    {
        printf("Empty priority queue!\n");
        exit(0);
    }
    return head->priority;
}
```

- Extract the element with maximum priority (minimum element)

```
struct node* pop(struct node* head){
    if(isEmpty(head))
    {
        printf("Empty priority queue!\n");
        exit(0);
    }
    struct node* temp = head;
    head = head->next;
    free(temp);
    // Return new head
    return head;
}
```



List Implementations

- Push new element at the right place

```
struct node* push(struct node* head, int newNum){
    // Create new Node
    struct node* nNode = newNode(newNum);
    // The head of list has lesser priority than new node
    if(head->priority>newNum)
    {
        nNode->next = head;
        // Return new head
        return nNode;
    }
    else
    {
        // Traverse the list and find the position
        // to insert the new node
        struct node* start = head;
        struct node* startPrev = head;
        while(newNum>start->priority)
        {
            if(start->next!=NULL)
            {
                startPrev = start;
                start = start->next;
            }
            else
            {
                // It has reached the end of list. Insert the new node
                start->next = nNode;
                // Return head
                return head;
            }
        }
        // Store what is now pointed by 'next' field of the node whose
        // priority is just above it (previous node of current start)
        nNode->next = startPrev->next;
        // Connect 'next' field of previous node to newNode's next
        startPrev->next = nNode;

        // Return head
        return head;
    }
}
```




List Implementations

- Complexities?
 - `getMin()` - $O(1)$
 - `pop()` - $O(1)$
 - `push()` - $O(n)$



Binary Search Tree Implementations


- A priority queue can also be entirely mapped onto a BST implementation
- `getMin()`
 - Use `getMin()` - $O(\log n)$
- `pop()`
 - Use `delete(getMin())` - $O(\log n)$
- `push()`
 - Use `insert()` method of BST - $O(\log n)$
- Conclusion on naïve implementations
 - A list implementation gives `getMin()/pop()` in $O(1)$
 - But `push()` in $O(n)$
 - A BST implementation gives `push()` in $O(\log n)$
 - But `getMin()/pop()` in $O(\log n)$ as well
 - How to get the best of both worlds?





Binary Heap



- Most common data structure for implementing a priority queue
- So ubiquitous in implementing priority queues that the word '*heap*' is used without any qualifier in this context
- A heap is a **binary tree** with two properties
 - Heap-order property
 - Structure property
- Operations on heaps can destroy one or more of these properties
- So a heap operation must not terminate until all heap properties are in order





Heap-Order Property

- Since we want to find the minimum quickly, it makes sense that the smallest element should be at the root.
- Continuing, any node should be smaller than all of its descendants
- Giving two types of Heaps
 - **min-heaps**: each node's key is less than or equal to each of its children
 - **max-heaps**: each node's key is greater than or equal to each of its children




- By design, item with **highest priority** is always the root
- `getMin()` is $O(1)$




Structure Property

- A heap is a complete binary tree
 - All levels are completely filled, apart from possibly the last
 - The last level is packed to the left
- Guarantee of height in $O(\log n)$



- Easy representation

- Complete binary tree is so regular, it can be represented in an array
- No need for complicated link management, and fast traversal



- Root is always at index 1

For node at index i^* :

- Left child at index $2i$
- Right child at index $2i + 1$
- Parent (if not root) at index $\lfloor \frac{i}{2} \rfloor$

*Heap could start at array index 0, but node indexing would become a little more complex



Implementation

- Min-heap version
- Simple array (capacity incremented by 1 for the unused first element)

```
#define CAPACITY 10  
  
// Global variable to track the size of Heap  
int size = 0;
```

```
// Remember we are starting at index 1  
int H[CAPACITY+1];
```

```
// Helper methods for indices  
int Root(){  
    return 1;  
}  
  
int Parent(int n){  
    return n/2;  
}  
  
int LeftChild(int n){  
    return 2*n;  
}  
  
int RightChild(int n){  
    return 2*n + 1;  
}
```


```
// Helper methods for node testing  
bool HasParent(int n){  
    return n!=Root();  
}  
  
bool IsNode(int n){  
    return n<=size;  
}
```

Source: UC Davis, ECS 36C course, Spring 2020




Implementation

- `getMin()` – Returns the node with highest priority (minimum value in min-heap)
- It is as simple as returning the root only



```
// Function to get the node with highest
// priority minimum element (root)
int getMin(int H[])
{
    if(size == 0)
    {
        printf("Empty priority queue!\n");
        exit(0);
    }
    return H[Root()];
}
```




- Next, we will insert or `push()` an element into the heap
- What is the most natural position for the new element? (Don't worry about the heap order property, just keep the structure property intact)




Implementation


- **push()** – Insert only at locations that keeps the tree complete



- Continue going up tree until heap-order is respected



- If heap-order not broken, stop!
- Otherwise, swap with parent




- Strategy known as shift up
- Also called bubble up, heapify-up

```
// Function to shift up the node in order
// to maintain the heap property
void shiftUp(int H[], int n){
    while (HasParent(n) && (H[Parent(n)] > H[n]))
    {
        swap(&H[Parent(n)], &H[n]);
        // Inside the while loop change
        // n to go to its parent
        n = Parent(n);
    }
}
```



Implementation

- Continue going up tree until heap-order is respected



- Strategy known as shift up
- Also called bubble up, heapify-up

```
// Function to shift up the node in order
// to maintain the heap property
void shiftUp(int H[], int n){
    while (HasParent(n) && (H[Parent(n)] > H[n]))
    {
        swap(&H[Parent(n)], &H[n]);
        // Inside the while loop change
        // n to go to its parent
        n = Parent(n);
    }
}
```

- Insertion of new item


```
void push(int H[], int newNum){
    // Check if heap is full
    if(size == CAPACITY)
    {
        printf("Priority queue full!\n");
        exit(0);
    }
    // Insert at the end
    H[size+1] = newNum;
    size++;
    // Shift up
    shiftUp(H, size);
}
```

Source: UC Davis, ECS 36C course, Spring 2020




Implementation


- `pop()` – Removing min item leaves hole at root. Move last item to root to keep the tree complete



- Continue going down tree until heap-order is respected again



- If heap-order not broken, stop!
- Otherwise, swap with smallest child



- Strategy known as shift down
- Also called bubble down, heapify-down



Implementation

- Strategy known as shift down
- Also called bubble down, heapify-down

```
// Function to shift down the node in order
// to maintain the heap property
void shiftDown(int H[], int n){
    // While node has at least one child
    // (if one, necessarily on the left)
    while(IsNode(LeftChild(n)))
    {
        // Consider left child by default
        int child = LeftChild(n);
        // If right child exists and smaller than
        // left child, then consider right child
        if(IsNode(RightChild(n)) && (H[RightChild(n)] < H[LeftChild(n)]))
            child = RightChild(n);
        // Exchange smallest child with node to
        // restore heap-order if necessary
        if(H[n]>H[child])
            swap(&H[n], &H[child]);
        else
            break;

        // Inside the while loop change
        // n to go one level down
        n = child;
    }
}
```

```
// Function to extract the element with
// maximum priority (in min-heap, this is
// node with the min value)
void Pop(int H[]){
    if(size == 0)
    {
        printf("Empty priority queue!\n");
        exit(0);
    }
    // Move last item back to root and reduce
    // heap's size
    H[Root()] = H[size];
    size--;
    shiftDown(H, Root());
}
```



Conclusion

- Running time complexities

Push	Pop	Top
$O(\log n)$	$O(\log n)$	$O(1)$

- Other operations:

- In-place item modification, e.g., when heap/priority queue is used for process scheduling, change priority of a process
- DecreaseItem(Item, DeltaPriority)
- IncreaseItem(Item, DeltaPriority)
- Heap variants:
 - 2-3 heap, Binomial heaps, d-ary heaps, Fibonacci heaps, Leftist heaps, Skew heaps etc.

Type	Push	Pop	Top
Sorted list	$O(n)$	$O(1)$	$O(1)$
BST	$O(\log n)$	$O(\log n)$	$O(\log n)$
Binary heap	$O(\log n)$	$O(\log n)$	$O(1)$



Heapify/BuildHeap



- Sometimes, build an entire heap directly out of an initial collection of items
- One use-case is Heapsort – an efficient way to sort an array
- Naïve approach (also known as Williams' method)
 - For all N items, `push()` N times
- Each `push()` operation takes $O(1)$ average time and $O(\log N)$ worst-case time
- For N items, this algorithm will run in $O(N)$ average time and $O(N \log N)$ worst-case time
- Is there a better way?



Heapify/BuildHeap


- Linear approach (Floyd's method)
 - Put all the elements on binary tree, only respecting structure property
 - Shift down all the elements starting from the first node who has at least one child, and up to the root

```
// Array  
int arr[] = {51, 43, 93, 18, 42, 99, 54, 2, 74};
```





Heapify/BuildHeap





Heapify/BuildHeap Implementations

```
// Main code
int main()
{
    // Array
    int arr[] = {51, 43, 93, 18, 42, 99, 54, 2, 74};
    // Size of array
    int n = sizeof(arr)/sizeof(arr[0]);
    // Print the array
    printf("Array: ");
    for (int i = 0; i < n; i++)
        printf("%d ", arr[i]);
    printf("\n");

    // Remember we are starting at index 1
    int H[CAPACITY+1];
    // Building heap
    buildheap(arr, H, n);
    // Print after heapifying
    printf("Min-Heap: ");
    printHeap(H, size);

    return 0;
}
```


```
void buildheap(int arr[], int H[], int n){
    // Only enforce structure property
    for(int i = 0; i < n; i++)
        H[i+1] = arr[i];
    // Set the global variable size's value
    size = n;

    // Now enforce heap-property
    for (int i = size/2; i >= 1; i--)
        shiftDown(H, i);
}
```



Heapify/BuildHeap Complexity Analysis

- The running time of heapify operation is $\Theta(n)$ where n is the number of elements in the heap
- RECAP - Height:** The height of a node is the length of the longest path from a leaf to this node



- All leaves are at height 0
- The height of a tree is equal to the height of the root from the deepest leaf (which is always equal to the depth of the tree)




Heapify/BuildHeap Complexity Analysis

- The running time of heapify operation is $\Theta(n)$ where n is the number of elements in the heap
- Some intuitions first [Great resource on this topic - <https://stackoverflow.com/a/18742428>]
- The basic idea is after creating a complete binary tree - move an offending node until it satisfies the heap property
 - shiftUp - swaps a node that is less than its parent (thereby moving it up) until it is no smaller than the node above it
 - shiftDown - swaps a node that is more than its smallest child (thereby moving it down) until it is no larger than both nodes below it
- The number of operations required for shiftUp and shiftDown is proportional to the distance the node may have to move
- For shiftUp, it is the distance to the top of the tree, so shiftUp is expensive for nodes at the bottom of the tree
- For shiftDown, it is the distance to the bottom of the tree, so shiftDown is expensive for nodes at the top of the tree



Heapify/BuildHeap Complexity Analysis

- Although both operations are $O(\log n)$ in the worst case, only one node is at the top whereas half the nodes lie in the bottom layer
- So it shouldn't be too surprising that if we have to apply an operation to every node, we would prefer shiftDown over shiftUp



Max # of
swaps
height h

Max # of
nodes
 $1 = 2^0$

height $h - 1$

$2 = 2^1$

height $h - 2$

$4 = 2^2$


height $h - 3$
 $= 0$

$8 = 2^3$



Heapify/BuildHeap Complexity Analysis

- Total swaps (max) = $\sum_{i=0}^h 2^i (h - i)$
- $A = h + 2(h - 1) + 2^2(h - 2) + \dots + 2^{h-1}1 + 0 \dots (1)$
- $2A = 2h + 2^2(h - 1) + 2^3(h - 2) + \dots + 2^h1 + 0 \dots (2)$



Max # of
swaps
height h

Max # of
nodes
 $1 = 2^0$

height $h - 1$

$2 = 2^1$

height $h - 2$

$4 = 2^2$

height $h - 3$
 $= 0$

$8 = 2^3$



Heapify/BuildHeap Complexity Analysis

- Total swaps (max) = $\sum_{i=0}^h 2^i (h - i)$
- $A = h + 2(h - 1) + 2^2(h - 2) + \dots + 2^{h-1}1 + 0 \dots (1)$
- $2A = 2h + 2^2(h - 1) + 2^3(h - 2) + \dots + 2^h1 + 0 \dots (2)$
- $(2) - (1) \Rightarrow A = -h + 2 + 2^2 + 2^3 + \dots + 2^{h-1} + 2^h$
 $= -h + \left(\frac{2^{h+1} - 1}{2 - 1} - 1 \right) = (2^{h+1} - 1) - (h + 1)$
 $= (N - 1) - (\log N + 1) = O(N)$







Heapsort

- Heapsort uses heaps to sort in $O(N \log N)$ time
- Basic strategy:
 - Build a min-heap of N elements in $O(N)$ time
 - Perform N `pop()` operations
 - Store these elements in a second array giving N sorted elements
- Since each `pop()` takes $O(\log n)$ times, total running time is $O(N \log N)$
- The main problem with this algorithm is that it uses an extra array
- A clever way to avoid this is to use the fact that after each `pop()` operation, the heap shrinks by 1
- Thus, the previous last cell can be used to store the element that was just popped
- Using this strategy, the array will contain the elements in decreasing sorted order
- If we want in more typical increasing sorted order, we can change the heap to `max-heap` so that parent has larger element than child







Heapsort










Heapsort






Heapsort







Heapsort




Swap



0	1	2	3	4	5	6
	6	5	1	9	10	12


0	1	2	3	4	5	6
	1	5	6	9	10	12



ShiftDown



0	1	2	3	4	5	6
	5	1	6	9	10	12





Remove

0	1	2	3	4	5	6
	1	5	6	9	10	12




Heapsort



0	1	2	3	4	5	6
	5	1	6	9	10	12

0	1	2	3	4	5	6
	1	5	6	9	10	12



0	1	2	3	4	5	6
	1	5	6	9	10	12

0	1	2	3	4	5	6
	1	5	6	9	10	12



Thank You